

DATA 621 HW2

Ahsanul Choudhury

March 17, 2018

```
if (!require('ggplot2')) (install.packages('ggplot2'))
if (!require('caret')) (install.packages('caret'))
if (!require('pROC')) (install.packages('pROC'))
```

1. Download the classification output data set (attached in Blackboard to the assignment).

```
my_data <- read.csv('classification-output-data.csv', header=T)
head(my_data)
```

```
##   pregnant glucose diastolic skinfold insulin  bmi pedigree age class
## 1         7     124         70      33     215 25.5   0.161  37     0
## 2         2     122         76      27     200 35.9   0.483  26     0
## 3         3     107         62      13      48 22.9   0.678  23     1
## 4         1      91         64      24       0 29.2   0.192  21     0
## 5         4      83         86      19       0 29.3   0.317  34     0
## 6         1     100         74      12      46 19.5   0.149  28     0
##   scored.class scored.probability
## 1           0         0.32845226
## 2           0         0.27319044
## 3           0         0.10966039
## 4           0         0.05599835
## 5           0         0.10049072
## 6           0         0.05515460
```

2. The data set has three key columns we will use:

- class: the actual class for the observation
- scored.class: the predicted class for the observation (based on a threshold of 0.5)
- scored.probability: the predicted probability of success for the observation

Use the `table()` function to get the raw confusion matrix for this scored dataset. Make sure you understand the output. In particular, do the rows represent the actual or predicted class? The columns?

```
t <- table(my_data$scored.class, my_data$class)
knitr::kable(t)
```

	0	1
0	119	30
1	5	27

The rows represent the predicted class and the columns represent the actual class.

3. Write a function that takes the data set as a dataframe, with actual and predicted classifications identified, and returns the accuracy of the predictions.

$$Accuracy = \frac{TP + TN}{TP + FP + TN + FN}$$

Accuracy

```
a_fun <- function(df){  
  df <- as.data.frame(table(my_data$scored.class, my_data$class))  
  tp <- df$Freq[4]; fp <- df$Freq[2]; fn <- df$Freq[3]; tn <- df$Freq[1]  
  accuracy <- ((tp + tn)/(tp + fp + tn + fn))  
  return (accuracy)  
}  
a_fun(df)
```

```
## [1] 0.8066298
```

4. Write a function that takes the data set as a dataframe, with actual and predicted classifications identified, and returns the classification error rate of the predictions.

$$\text{Classification Error Rate} = \frac{FP + FN}{TP + FP + TN + FN}$$

Verify that you get an accuracy and an error rate that sums to one.

Classification Error Rate

```
err_fun <- function(df){  
  df <- as.data.frame(table(my_data$scored.class, my_data$class))  
  tp <- df$Freq[4]; fp <- df$Freq[2]; fn <- df$Freq[3]; tn <- df$Freq[1]  
  error <- ((fp + fn)/(tp + fp + tn + fn))  
  return (error)  
}  
err_fun(df)
```

```
## [1] 0.1933702
```

Verify

Accuracy + Classification Error Rate = 1

5. Write a function that takes the data set as a dataframe, with actual and predicted classifications identified, and returns the precision of the predictions.

$$\text{Precision} = \frac{TP}{TP + FP}$$

Precision

```
prec_fun <- function(df){  
  df <- as.data.frame(table(my_data$scored.class, my_data$class))  
  tp <- df$Freq[4]; fp <- df$Freq[2]  
  prec <- tp/(tp + fp)  
  return (prec)  
}  
prec_fun(df)
```

```
## [1] 0.84375
```

6. Write a function that takes the data set as a dataframe, with actual and predicted classifications identified, and returns the sensitivity of the predictions. Sensitivity is also known as recall.

$$\text{Sensitivity} = \frac{TP}{TP + FN}$$

Sensitivity

```
sen_fun <- function(df){  
  df <- as.data.frame(table(my_data$scored.class, my_data$class))  
  tp <- df$Freq[4]; fn <- df$Freq[3]  
  sen <- tp/(tp + fn)  
  return (sen)  
}  
sen_fun(df)  
  
## [1] 0.4736842
```

7. Write a function that takes the data set as a dataframe, with actual and predicted classifications identified, and returns the specificity of the predictions.

$$Specificity = \frac{TN}{TN + FP}$$

Specificity

```
spec_fun <- function(df){  
  df <- as.data.frame(table(my_data$scored.class, my_data$class))  
  fp <- df$Freq[2]; tn <- df$Freq[1]  
  spec <- tn/(tn + fp)  
  return (spec)  
}  
spec_fun(df)  
  
## [1] 0.9596774
```

8. Write a function that takes the data set as a dataframe, with actual and predicted classifications identified, and returns the F1 score of the predictions.

$$F1\ Score = \frac{2 \times Precision \times Sensitivity}{Precision + Sensitivity}$$

F1 Score

```
f1_fun <- function(df){  
  df <- as.data.frame(table(my_data$scored.class, my_data$class))  
  tp <- df$Freq[4]; fp <- df$Freq[2]; fn <- df$Freq[3]  
  prec <- tp/(tp + fp); sen <- tp/(tp + fn)  
  f1 <- (2 * prec * sen) / (prec + sen)  
  return (f1)  
}  
f1_fun(df)  
  
## [1] 0.6067416
```

9. Before we move on, let's consider a question that was asked: What are the bounds on the F1 score? Show that the F1 score will always be between 0 and 1. (Hint: If $0 < a < 1$ and $0 < b < 1$ then $ab < a$).

Since F1 score is a function of precision and sensitivity which are themselves bounded between 0 and 1, the solution to equation, $F1\ Score = \frac{2 \times Precision \times Sensitivity}{Precision + Sensitivity}$ will always be between 0 and 1.

e.g.

```

set.seed(1023)
prec <- runif(50, 0, 1)
sen <- runif(50, 0, 1)
f1 <- (2 * prec * sen) / (prec + sen)
result <- summary(f1)
c(result[1], result[6])

```

```

##      Min.      Max.
## 0.03142 0.92510

```

10. Write a function that generates an ROC curve from a data set with a true classification column (class in our example) and a probability column (scored.probability in our example). Your function should return a list that includes the plot of the ROC curve and a vector that contains the calculated area under the curve (AUC). Note that I recommend using a sequence of thresholds ranging from 0 to 1 at 0.01 intervals.

```

my_roc <- function(class, results){
  class <- class[order(results, decreasing=TRUE)]
  df <- data.frame(TPR=(cumsum(class)/sum(class)), FPR=(cumsum(!class)/sum(!class)), class)

  dFPR <- c(diff(df$FPR), 0)
  dTPR <- c(diff(df$TPR), 0)
  AUC <- sum(df$TPR * dFPR) + sum(dTPR * dFPR) / 2


  results <- list(df, AUC)
  return(results)
}

my_roc_auc <- my_roc(my_data$class, my_data$scored.probability)
my_roc_results <- my_roc_auc[[1]]
auc <- my_roc_auc[[2]]

ggplot(my_roc_results, aes(FPR, TPR)) +
  geom_line(color='red') +
  labs(title = "ROC Curve Using Function" , x = "Specificity", y = "Sensitivity") +
  geom_abline(linetype=2)

```

ROC Curve Using Function



```
auc
```

```
## [1] 0.8503113
```

```
# reference: http://blog.revolutionanalytics.com/2016/08/roc-curves-in-two-lines-of-code.html
# http://blog.revolutionanalytics.com/2016/11/calculating-auc.html
```

11. Use your **created R functions** and the provided classification output data set to produce all of the classification metrics discussed above.

```
Classification <- c('Accuracy', 'Classification Error Rate', 'Precision',
                   'Sensitivity', 'Specificity', 'F1 Score')
Results <- round(c(a_fun(df), err_fun(data), prec_fun(df), sen_fun(df),
                  spec_fun(df), f1_fun(df)), 4)
df1 <- as.data.frame(cbind(Classification, Results))
knitr::kable(df1)
```

Classification	Results
Accuracy	0.8066
Classification Error Rate	0.1934
Precision	0.8438
Sensitivity	0.4737
Specificity	0.9597
F1 Score	0.6067

12. Investigate the **caret** package. In particular, consider the functions `confusionMatrix`, `sensitivity`, and `specificity`. Apply the functions to the data set. How do the results compare with your own functions?


```
confusionMatrix(my_data$scored.class, my_data$class, positive="1")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0    1
##           0 119  30
##           1   5  27
##
##           Accuracy : 0.8066
##           95% CI : (0.7415, 0.8615)
##           No Information Rate : 0.6851
##           P-Value [Acc > NIR] : 0.0001712
##
##           Kappa : 0.4916
##           McNemar's Test P-Value : 4.976e-05
##
##           Sensitivity : 0.4737
##           Specificity : 0.9597
##           Pos Pred Value : 0.8438
##           Neg Pred Value : 0.7987
##           Prevalence : 0.3149
##           Detection Rate : 0.1492
##           Detection Prevalence : 0.1768
##           Balanced Accuracy : 0.7167
##
##           'Positive' Class : 1
##
```

The caret package result and my results are the same.

13. Investigate the **pROC** package. Use it to generate an ROC curve for the data set. How do the results compare with your own functions?

```
r_curve <- roc(my_data$class~my_data$scored.probability)
plot(r_curve, main= "ROC Curve using pROC Package")
```



```
auc(roc(my_data$class, my_data$scored.probability))
```

```
## Area under the curve: 0.8503
```

The results are exact match with my functions.